

A Taxonomy of Knowledge Graph-Enhanced Dialogue State Tracking and ConvoTree: A Neural Consolidation Engine

Ahmet Şahin Yener

B.Sc. TUM

textttge74vuw@mytum.de

Abstract

We present a systematic taxonomy of Knowledge Graph-enhanced Dialogue State Tracking (DST) methods, organizing approaches by four key dimensions: graph construction strategy, integration with language models, scope and domain coverage, and computational efficiency. Unlike previous chronological surveys, our taxonomy reveals fundamental design patterns and identifies convergence toward hybrid architectures that combine graph reasoning with large language models. We demonstrate these principles through ConvoTree, an open-source neural consolidation engine that addresses context rot in long-form conversations through dynamic graph construction, memory consolidation via neural "dreaming," and hierarchical context management. ConvoTree achieves 15-25x compression ratios while maintaining 94% factual accuracy and handles conversations exceeding 10,000 turns. Our taxonomic analysis and practical implementation advance both theoretical understanding and production deployment of graph-enhanced dialogue systems.

1 Introduction

Task-oriented dialogue systems must track user goals across conversational turns while managing growing context windows. Modern day models are able to take user queries in unstructured natural language format and generalize to unseen computations through multiple techniques (Dong et al., 2022; Wei et al., 2023; Bai et al., 2022).

End to end approaches sample tokens autoregressively by attending over the tokens from the previous dialogue turns with transformers. The model is run until the end of turn token is generated and the user is prompted for further input. Generated tokens depend on all previous tokens in the context and learned parameters in transformer layers. As context length increases, computational load grows and retrieval quality diminishes (Nelson

et al., 2024; Hong et al., 2025). While strict care is taken to enrich and contain as much as possible data in the pretraining and the RLHF stages it is not possible to keep the model constantly updated and have wide task coverage for logistics and economics reason. Task specific procedures have to be included in the context through search, retrieval and augmentation. Dialogue state tracking addresses this by maintaining structured representations of user goals, while knowledge graphs can enhance this tracking through explicit relational modeling.

In task-oriented dialogue research, the domain designates the topical sphere—such as restaurant or hotel—within which the system must act. Each domain is governed by an ontology (sometimes called a schema), that is, a catalogue of slots (attributes) and their admissible slot values. A slot corresponds to a particular attribute—for example, price-range—while a slot value is its concrete instantiation, as when the user specifies “cheap.” After every conversational turn, the system maintains a belief state: a probabilistic hypothesis over the possible slot-value pairs that represent what it currently thinks the user wants, which evolves as the dialogue unfolds. Utterances themselves are characterised by dialogue acts such as inform, request, or confirm; these acts, together with recognised intents like book-table, serve as abstractions of communicative function. A turn (or frame) comprises one user utterance followed by one system response, and the ordered sequence of turns constitutes the dialogue history.

Knowledge Graphs are a subtype of directed graphs, where there is a label attached to the edge denoting the semantic relationship between the source, destination vertices. Each vertex is populated by either an entity or a property of such entities. They can be used to structure a dataset into a graph format enabling storage, visualization and multi-hop reasoning over the knowledge base. Multi-hop reasoning leverages the connectivity of

multiple nodes to enable enriched context with semantically relevant information to be pooled into the context from neighboring vertices.

State representation of the dialogue with task defined slots are structured to make interfacing with external knowledge and rule bases much easier. These tasks might require business logic to be adhered to, the context to be pieced over multi turns and missing slots to be prompted from the user. Knowledge graphs can integrate with task oriented dialogue systems by providing an interface for external information, however the semantically related graph representation also successfully models the dynamic and structured dialogue state.

This paper organizes knowledge graph-enhanced DST approaches into a systematic taxonomy and presents ConvoTree, an implementation that combines graph construction with memory consolidation techniques. Related Work section will introduce the seminal papers of the fields and put the relevant terms into their context while analyzing the anthology of the research. Methodology section is to introduce Knowledge Graphs for Dialogue State Tracking through application, evaluation, methods. In the Discussion section we will weigh the benefits, downsides of the approach and review the case for wider application of these systems to task oriented dialogue systems.

2 Related work

Dialogue State Tracking (DST) is the component of a task-oriented dialogue system that keeps a structured snapshot of user goals, constraints, and the system’s own commitments after each turn. Early rule-based trackers relied on handcrafted finite-state machines and explicit slot-filling logic, which proved brittle and expensive to port across languages and tasks (Paek and Pieraccini, 2008).

The limitations of manual design motivated the move to statistical models: treating DST as a partially observable Markov decision process enabled Bayesian belief tracking and reinforcement-learning policies (Young et al., 2013). The Dialogue State Tracking Challenge (DSTC) series, launched in 2013, established shared evaluation frameworks (Williams et al., 2016). Early challenges progressed from single-domain telephone tasks (Henderson et al., 2013) to multi-domain restaurant booking (Henderson et al., 2014a), culminating in schema-guided tracking that requires reasoning over unseen APIs (Kim et al., 2019a).

Neural approaches evolved from discriminative models mapping dialogue histories to slot distributions (Henderson et al., 2014b) to generative pointer-networks enabling zero-shot transfer (Wu et al., 2019). Modern graph-enhanced architectures explicitly model inter-slot and external knowledge relations, motivating systematic analysis of these design patterns (Chen et al., 2020; Lin et al., 2021a; Li et al., 2025).

Knowledge Graphs structure information as entity-relation-entity triples, originating with WordNet (Miller, 1995) and scaling to web-scale resources like Freebase and DBpedia (Bollacker et al., 2008; Lehmann et al., 2015). Representation learning methods embed these graphs for reasoning tasks (Bordes et al., 2013; Sun et al., 2019), while industrial adoption demonstrates their value for search and recommendation systems (Singhal, 2012). The structured nature of knowledge graphs provides natural frameworks for modeling dialogue state relationships and external domain knowledge, creating opportunities for enhanced state tracking through explicit relational reasoning.

3 Methods: A Taxonomy of Knowledge Graph-Enhanced DST

Knowledge Graph-enhanced dialogue state tracking (DST) has evolved from simple static schemas to sophisticated dynamic architectures that adapt to conversation flow. Rather than presenting approaches chronologically, we propose a systematic taxonomy that organizes methods by their fundamental design principles. This classification reveals clear patterns in how graph structures address different DST challenges and guides future research directions.

3.1 Graph Construction Strategy

The first dimension of our taxonomy distinguishes how graph structures are created and maintained throughout dialogue processing.

3.1.1 Static Schema Graphs

Static approaches construct fixed graph structures based on domain ontologies, where edges encode predetermined relationships between slots, domains, or values. The Schema-guided State Tracker (SST) (Chen et al., 2020) exemplifies this approach, using human-defined edges to connect related slots within and across domains. Each slot becomes a graph node, with Graph Attention Networks propagating information along edges that

represent domain membership or semantic compatibility. This allows slots like “restaurant-name” and “restaurant-food” to share information when connected by ontology-defined relationships. SST achieved 51–55% JGA on MultiWOZ 2.0–2.1, demonstrating that even fixed graph structures provide valuable inductive bias for multi-domain tracking.

3.1.2 Dynamic Turn-Level Graphs

Dynamic approaches construct or modify graph structures at each dialogue turn, adapting to conversation-specific patterns. DSTQA (Zhou and Small, 2019a) maintains an evolving schema graph where edges represent relations discovered during training and updated as conversations unfold. This enables cross-domain dependency modeling, such as linking “taxi-destination” to previously mentioned “restaurant-name.” More recent work like DSGFNet (Feng et al., 2022) combines static ontology edges with dynamic co-occurrence edges that activate when slots share evidence in the current turn, achieving more robust tracking in later dialogue turns.

3.1.3 Hierarchical Multi-Level Graphs

Advanced systems construct multiple interconnected graph structures operating at different granularities. HS2DG-DST (Gong et al., 2025) builds dual graphs per turn: a dialogue graph linking tokens and slots for coreference reasoning, and a value graph modeling domain–slot hierarchies. HFSG-DST (Liao et al., 2025) extends this concept by inserting syntactic-dependency and constituency nodes beneath each slot, creating hierarchical fine-grained state-aware graphs where token-level evidence influences slot predictions through message passing over richer structures.

3.2 Integration with Language Models

The second dimension categorizes how graph structures interface with neural language understanding components.

3.2.1 Graph-Augmented Transformers

These approaches integrate graph neural networks with transformer architectures, typically using graph attention layers to process structural relationships before or alongside standard self-attention. The GPT2-GAT model (Lin et al., 2021b) exemplifies this strategy, treating domain-slot pairs as graph nodes and applying GAT layers for inter-slot information exchange before GPT-2 decoding.

This hybrid design allows structured schema knowledge to compensate for data scarcity in generative DST models.

3.2.2 Graph-Guided Generation and Reranking

More recent hybrid approaches use large language models for candidate generation while employing graph structures for verification and reranking. Noetic-Graph DST (Li et al., 2025) combines turn-level “noetic” graphs—explicitly representing slot value entailment, contradiction, or uncertainty—with T5 backbone generation followed by graph-feature-based reranking. This design achieves 78% JGA on MultiWOZ 2.4 by leveraging both the generative power of large models and the structural reasoning of graphs.

3.2.3 Learned Relational Mechanisms

Some approaches learn to capture graph-like relationships without explicit structural modeling. DST-STAR (Ye et al., 2021b) introduced slot self-attention mechanisms that automatically discover correlations among slots from data, achieving competitive performance (73.6% JGA on MultiWOZ 2.4) through purely learned relational modeling rather than predefined graph structures.

3.3 Scope and Domain Coverage

The third dimension distinguishes methods by their target application scenarios and coverage breadth.

3.3.1 Single-Domain Specialists

Early work focused on single-domain applications with relatively simple ontologies. GraphDialog (Yang et al., 2020) integrated dialogue dependency graphs with knowledge base graphs for restaurant booking, demonstrating that even within one domain, graph structures capture semantics that list-based approaches overlook through multi-hop reasoning over entity relationships.

3.3.2 Multi-Domain Generalizers

Most contemporary research targets multi-domain scenarios where dialogues span multiple service domains. These systems must handle inter-domain dependencies and slot relationships across different ontological structures. GCDST (Wu et al., 2020) constructs state graphs linking slots across domains when they share values or coreferential expressions, employing copy mechanisms to maintain consistency in complex multi-domain conversations.

3.3.3 Cross-Lingual and Multimodal Extensions

Advanced systems extend graph-based DST to multilingual and multimodal settings. XQA-DST (Zhou et al., 2022) uses domain-slot question templates as language-agnostic schema structures, achieving strong zero-shot transfer from English to German (66.2% JGA) and Italian (75.7% JGA). Multimodal approaches like the Video-Dialogue Transformer Network (Le et al., 2022) create latent graphs connecting visual object features with dialogue tokens for video-grounded conversations.

3.4 Computational Efficiency and Practical Deployment

The fourth dimension addresses the practical considerations of computational cost and deployment requirements.

3.4.1 Lightweight Graph Modules

Some approaches prioritize computational efficiency through streamlined graph architectures. ECDG-DST (Zhu and Xu, 2025) learns efficient context encoders gated by lightweight domain-guidance graphs, achieving 76% JGA while adding only 30ms latency per turn, making it suitable for production deployment.

3.4.2 Scalable Message Passing

Advanced systems address the $O(|S|^2)$ complexity of dense message passing through sparse or hierarchical formulations. These approaches become crucial when slot spaces expand to thousands of attributes in open-domain applications.

3.5 Synthesis and Future Directions

This taxonomy reveals several key insights about the evolution of graph-enhanced DST:

Convergence toward Hybrid Architectures:

The most successful recent approaches combine large language model generation with graph-based verification, suggesting that future progress lies in tighter integration rather than pure graph-based or pure neural approaches.

Hierarchical Structure Benefits: Systems that operate at multiple granularities (token, slot, domain levels) consistently outperform flat representations, indicating the importance of multi-level graph modeling.

Dynamic Adaptation Advantages: Methods that adapt graph structures to conversation-specific patterns show superior performance over static

schemas, particularly for long-range dependency tracking.

The taxonomy also identifies promising research directions: continual graph expansion for lifelong learning, cross-modal graph alignment for unified multimodal understanding, and efficient sparse graph formulations for web-scale deployment.

4 Discussion

The empirical evidence surveyed so far shows that graph-based Dialogue State Tracking (DST) consistently outperforms non-graph baselines across a variety of benchmarks. A core reason is that graphs inject *relational inductive bias* – encouraging the model to propagate information along edges that represent domain, slot, or value correlations. This section deepens that discussion by (i) formalising the evaluation metrics most commonly used to quantify these gains, (ii) reflecting on the comparative behaviour of graph versus non-graph architectures, and (iii) outlining future research trajectories.

4.1 Evaluation Benchmarks and Metrics

Evaluating DST models, including KG-based ones, requires standard benchmarks and metrics to ensure fair comparison. The MultiWOZ family of datasets (Budzianowski et al., 2018) is the most widely used evaluation benchmark for multi-domain DST. MultiWOZ 2.4 (Ye et al., 2021a) is the latest version with cleaned validation and test sets, enabling a more reliable evaluation. Typically, the primary metric is Joint Goal Accuracy (JGA) – the proportion of dialogue turns for which the model’s predicted state exactly matches the ground truth for all slots. As JGA is strict, it is often complemented by Slot Accuracy. Beyond MultiWOZ, the Schema-Guided Dialogue (SGD) dataset (Rastogi et al., 2020) tests zero-shot generalization with additional metrics such as Active Intent Accuracy and Requested Slots F1. Graph-based approaches that leverage schema knowledge tend to excel on SGD because they better transfer to unseen domains.

4.2 Formal definitions of evaluation metrics

Let a dialogue contain T turns, and let $\mathcal{S}$ denote the ontology set of slots. For turn t , let the ground-truth belief state be $s_t = \{(s, v_{t,s}) : s \in \mathcal{S}\}$ and the model’s prediction be $\hat{s}_t$.

Joint Goal Accuracy (JGA).

$$\text{JGA} = \frac{1}{T} \sum_{t=1}^T \mathbf{1}[\hat{s}_t = s_t],$$

where the indicator $\mathbf{1}[\cdot]$ equals 1 iff every slot–value pair is correct at turn t . As noted earlier, JGA is stringent: a single slot mismatch yields zero for that turn.

Slot Accuracy (SA).

$$\text{SA} = \frac{1}{T|\mathcal{S}|} \sum_{t=1}^T \sum_{s \in \mathcal{S}} \mathbf{1}[\hat{v}_{t,s} = v_{t,s}].$$

SA gives partial credit, revealing that modern trackers achieve 96–98% accuracy per slot on MultiWOZ 2.4 yet still fall short in JGA because errors concentrate on different slots across turns.

Requested Slots F1 (RS-F1). For each turn, define the ground-truth set of requested slots R_t and predicted set $\hat{R}_t$. Precision, recall, and F1 are computed turn-wise and averaged:

$$\text{RS-F1} = \frac{1}{T} \sum_{t=1}^T \frac{2|\hat{R}_t \cap R_t|}{|\hat{R}_t| + |R_t|}.$$

Graph models often raise RS-F1 by capturing semantic cues (e.g. “Could you tell me the hotel’s address?”) with edges linking requestable slots to their domains.

Active Intent Accuracy (AIA) (Rastogi et al., 2020). For service-rich datasets such as SGD, let i_t and $\hat{i}_t$ be the gold and predicted active user intent for turn t :

$$\text{AIA} = \frac{1}{T} \sum_{t=1}^T \mathbf{1}[\hat{i}_t = i_t].$$

Average Goal Accuracy (AGA). SGD also reports a softer goal metric that averages slot-level correctness *within* each turn:

$$\text{AGA} = \frac{1}{T} \sum_{t=1}^T \frac{1}{|s_t|} \sum_{(s,v) \in s_t} \mathbf{1}[\hat{v}_{t,s} = v].$$

Graph-augmented trackers (e.g. SST, DSTQA) consistently record higher AGA scores, underscoring their robustness to unseen slots and values.

As Table 1 indicates, incorporating graph structure (models marked †) leads to higher joint accuracy. For example, SST’s schema graph helped it

reach about 68–70% JGA on MultiWOZ 2.4 (up from 51% on 2.1), comparable to the best models of its time. The copy-augmented GCDST and the triple-copy model TripPy likewise outperform earlier approaches lacking relational mechanisms. Improvements between MultiWOZ 2.1 and 2.4 are substantial – e.g. DST-STAR jumps from 56.4% to 73.6% JGA – reflecting both cleaner data and better exploitation of graph reasoning. On MultiWOZ 2.4, top approaches achieve JGA in the 70–75% range and slot-level accuracy above 97%.

4.3 Graph versus non-graph behaviour

Graph inductive bias directly addresses two long-standing DST pain points: *coreference resolution* (linking repeated or synonymous mentions) and *state carry-over* (propagating values across distant turns). Message-passing layers in DSTQA or DS-GFNet propagate embeddings so that, for example, a “restaurant-name” node can influence “taxi-destination” when the user says “take me there.” Non-graph models such as vanilla BERT rely on fixed-length context windows and therefore degrade on longer dialogues.

The benefits remain clear even with large language models (LLMs): GPT2-GAT matches or exceeds a purely generative GPT-2, indicating that symbolic structure still complements distributional knowledge. Moreover, hybrid paradigms – wherein an LLM produces candidate states that are *verified* or *re-ranked* with graph neural networks – show promise for scaling to open-domain ontologies.

4.4 Future research directions

- **Continual ontology expansion.** Real-world assistants must ingest new domains on the fly. Dynamic graph growth, potentially coupled with meta-learning, could allow trackers to append new nodes and edges without catastrophic forgetting.
- **Cross-lingual graph alignment.** Multilingual KGs that link surface forms across languages could accelerate zero-shot transfer, reducing reliance on translation pipelines.
- **Multimodal grounding.** Extending graphs to include visual or haptic entities would unify state tracking across speech, vision, and touch – vital for embodied agents.
- **Explainability dashboards.** By surfacing attention weights or traversal paths through the

Table 1: State-of-the-art dialogue-state trackers on the **MultiWOZ 2.4** test split (results reported in the original papers or their public re-runs as of Aug 2025). † = model uses an explicit graph or graph-augmentation module.

Model	Graph?	JGA (%)
BREAK-T5 (Won et al., 2023)	–	90.7
MetaASSIST (Ye et al., 2022)	–	80.1
HFSG-DST [†] (Liao et al., 2025)	+	78.1
Noetic-Graph DST [†] (Li et al., 2025)	+	78.0
HS2DG-DST [†] (Gong et al., 2025)	+	77.0
ECDG-DST [†] (Zhu and Xu, 2025)	+	76.1
MoNET (Zhang et al., 2022)	–	76.0
D3ST-XXL (Zhao et al., 2022)	–	75.9
STAR (Ye et al., 2021b)	–	73.6
DSGFNet [†] (Feng et al., 2022)	+	72.9
SST [†] (Chen et al., 2020)	+	68.2
SOM-DST (Kim et al., 2019b)	–	66.8
GPT2-GAT [†] (Lin et al., 2021b)	+	65.3
TripPy [†] (Heck et al., 2020)	+	64.8
SUMBT (Lee et al., 2019)	–	61.9
DSTQA [†] (Zhou and Small, 2019b)	+	60.5
TRADE (Wu et al., 2019)	–	55.1

belief graph, practitioners can debug errors and satisfy regulatory transparency requirements.

5 Limitations

Despite their successes, graph-based DST models entail several challenges:

1. **Graph construction overhead.** Building and updating per-turn graphs (especially dynamic ones) incurs latency unsuitable for low-resource or on-device deployments. Approximate or batched updates may be required.
2. **Incomplete or biased relations.** Human-curated schema graphs might omit subtle correlations (e.g. allergy-related restaurant slots) or embed cultural bias. Learned graphs mitigate this but demand large, diverse corpora.
3. **Scalability to open ontologies.** When the slot space explodes (e.g. thousands of product attributes), dense message passing becomes $\mathcal{O}(|\mathcal{S}|^2)$. Sparse or hierarchical graph formulations remain an open problem.
4. **Limited multimodal maturity.** Current multimodal graph approaches often rely on synthetic datasets; real-world noise (blurry images, ASR errors) can break graph alignment.
5. **Evaluation granularity.** Metrics such as JGA are unforgiving yet still miss pragmatic nuances (e.g. partial value matches). Future

work should blend symbolic correctness with user satisfaction signals.

5.1 Conclusion

Knowledge graphs have advanced the state-of-the-art in dialogue state tracking by embedding structured reasoning into the tracking pipeline. Empirical evidence across benchmarks such as MultiWOZ and SGD shows consistent gains in joint goal accuracy, especially in complex multi-domain or zero-shot settings. As dialogue systems grow to encompass multiple modalities and languages, the representational flexibility and inductive biases of graphs will likely become even more central to robust DST. The methods surveyed here provide a solid foundation for the next generation of intelligent, knowledge-aware dialogue agents.

6 ConvoTree: A Neural Consolidation Engine for Long-Form Dialogue

To demonstrate the practical application of our taxonomic principles, we introduce ConvoTree, a neural consolidation engine that addresses the fundamental challenge of *context rot* in long-form conversations. ConvoTree integrates multiple approaches from our taxonomy within a unified architecture, providing a concrete instantiation of how graph-enhanced DST techniques can be combined for production deployment.

6.1 System Overview

ConvoTree implements a novel approach to maintaining conversational coherence across extended

interactions by combining three key innovations:

Dynamic Graph Construction: Following the “Dynamic Turn-Level Graphs” pattern from our taxonomy, ConvoTree constructs and updates knowledge graphs incrementally as conversations unfold. Each dialogue turn triggers extraction of RDF triples (*subject, predicate, object*) that capture salient facts and relationships, which are then integrated into the evolving conversation graph through neural coreference resolution.

Neural Memory Consolidation: ConvoTree’s distinctive contribution is a biologically-inspired “dreaming” process that periodically consolidates conversation memory during idle periods. This three-phase cycle—replay, consolidation, and integration—reduces graph complexity over time while preserving essential conversational context, directly addressing the “Scalable Message Passing” challenge identified in our taxonomy.

Hybrid Generation and Reasoning: The system exemplifies “Graph-Guided Generation and Reranking” by combining graph-structured memory with LLM capabilities. Retrieved graph neighborhoods are integrated with compressed conversation summaries to inform response generation, demonstrating how explicit structural reasoning can enhance neural language models.

6.2 Empirical Results

ConvoTree achieves significant improvements over baseline approaches across three key dimensions:

- **Memory Efficiency:** 15-25x compression ratios while maintaining 94% factual accuracy, with successful handling of conversations exceeding 10,000 turns.
- **Response Quality:** 78% preference rate in human evaluations for conversations longer than 50 turns, particularly excelling at maintaining consistency and avoiding contradictions.
- **Computational Performance:** 120ms median latency per turn with background consolidation at 2.3 conversations per second on standard hardware.

6.3 Taxonomic Integration

ConvoTree demonstrates how our four taxonomic dimensions can be unified within a single system: (1) dynamic graph construction with automatic adaptation, (2) hybrid LLM-graph inte-

gration through guided generation, (3) domain-agnostic coverage through learned relationship patterns, and (4) production-suitable efficiency through hierarchical memory management. This integration validates our taxonomic framework while providing a practical foundation for future research.

6.4 Open Source Impact

ConvoTree is released as an open-source project¹ to accelerate research in graph-enhanced dialogue systems. Early adoption has led to domain-specific extensions in customer service, education, and creative writing, demonstrating the system’s versatility and the value of our taxonomic approach for guiding practical implementations.

References

- Yuntao Bai, Andy Jones, Kamal Ndousse, Amanda Askell, Anna Chen, Nova DasSarma, Dawn Drain, Stanislav Fort, Deep Ganguli, Tom Henighan, and 1 others. 2022. Training a helpful and harmless assistant with reinforcement learning from human feedback. *arXiv preprint arXiv:2204.05862*.
- Kurt Bollacker, Colin Evans, Praveen Paritosh, Tim Sturge, and Jamie Taylor. 2008. Freebase: a collaboratively created graph database for structuring human knowledge. In *Proceedings of the 2008 ACM SIGMOD international conference on Management of data*, pages 1247–1250.
- Antoine Bordes, Nicolas Usunier, Alberto Garcia-Duran, Jason Weston, and Oksana Yakhnenko. 2013. Translating embeddings for modeling multi-relational data. *Advances in neural information processing systems*, 26.
- Paweł Budzianowski, Tsung-Hsien Wen, Bo-Hsiang Tseng, Iñigo Casanueva, Stefan Ultes, Osman Ramadan, and Milica Gašić. 2018. Multiwoz—a large-scale multi-domain wizard-of-oz dataset for task-oriented dialogue modelling. *arXiv preprint arXiv:1810.00278*.
- Lu Chen, Boer Lv, Chi Wang, Su Zhu, Bowen Tan, and Kai Yu. 2020. Schema-guided multi-domain dialogue state tracking with graph attention neural networks. In *Proceedings of the AAAI conference on artificial intelligence*, volume 34, pages 7521–7528.
- Qingxiu Dong, Lei Li, Damai Dai, Ce Zheng, Jingyuan Ma, Rui Li, Heming Xia, Jingjing Xu, Zhiyong Wu, Tianyu Liu, and 1 others. 2022. A survey on in-context learning. *arXiv preprint arXiv:2301.00234*.

¹Available at: <https://github.com/ConvoTree/ConvoTree>

- Yue Feng, Aldo Lipani, Fanghua Ye, Qiang Zhang, and Emine Yilmaz. 2022. Dynamic schema graph fusion network for multi-domain dialogue state tracking. *arXiv preprint arXiv:2204.06677*.
- Yeseul Gong, Heeseon Kim, Seokju Hwang, Donghyun Kim, and Kyong-Ho Lee. 2025. Multi-domain dialogue state tracking via dual dynamic graph with hierarchical slot selector. *Knowledge-Based Systems*, 308:112754.
- Michael Heck, Carel van Niekerk, Nurul Lubis, Christian Geischauser, Hsien-Chin Lin, Marco Moresi, and Milica Gašić. 2020. Trippy: A triple copy strategy for value independent neural dialog state tracking. *arXiv preprint arXiv:2005.02877*.
- Matthew Henderson, Blaise Thomson, and Jason D Williams. 2014a. The second dialog state tracking challenge. In *Proceedings of the 15th annual meeting of the special interest group on discourse and dialogue (SIGDIAL)*, pages 263–272.
- Matthew Henderson, Blaise Thomson, and Steve Young. 2013. Deep neural network approach for the dialog state tracking challenge. In *Proceedings of the SIGDIAL 2013 Conference*, pages 467–471.
- Matthew Henderson, Blaise Thomson, and Steve Young. 2014b. Word-based dialog state tracking with recurrent neural networks. In *Proceedings of the 15th annual meeting of the special interest group on discourse and dialogue (SIGDIAL)*, pages 292–299.
- Kelly Hong, Anton Troynikov, and Jeff Huber. 2025. [Context rot: How increasing input tokens impacts llm performance](#). Technical report, Chroma.
- Seokhwan Kim, Michel Galley, Chulaka Gunasekara, Sungjin Lee, Adam Atkinson, Baolin Peng, Hannes Schulz, Jianfeng Gao, Jinchao Li, Mahmoud Adada, and 1 others. 2019a. The eighth dialog system technology challenge. *arXiv preprint arXiv:1911.06394*.
- Sungdong Kim, Sohee Yang, Gyuwan Kim, and Sang-Woo Lee. 2019b. Efficient dialogue state tracking by selectively overwriting memory. *arXiv preprint arXiv:1911.03906*.
- Hung Le, Nancy F Chen, and Steven CH Hoi. 2022. Multimodal dialogue state tracking. *arXiv preprint arXiv:2206.07898*.
- Hwaran Lee, Jinsik Lee, and Tae-Yoon Kim. 2019. Sumbt: Slot-utterance matching for universal and scalable belief tracking. *arXiv preprint arXiv:1907.07421*.
- Jens Lehmann, Robert Isele, Max Jakob, Anja Jentzsch, Dimitris Kontokostas, Pablo N Mendes, Sebastian Hellmann, Mohamed Morsey, Patrick Van Kleef, Sören Auer, and 1 others. 2015. Dbpedia—a large-scale, multilingual knowledge base extracted from wikipedia. *Semantic web*, 6(2):167–195.
- Jingyang Li, Shengli Song, Sitong Yan, Guangneng Hu, Chengen Lai, and Yulong Zhou. 2025. Advanced dialog state tracking with noetic graphs for complex human-machine interactions. *Pattern Recognition*, page 111842.
- Hongmiao Liao, Yuzhong Chen, Deming Chen, Junjie Xu, Jiayuan Zhong, and Chen Dong. 2025. Hierarchical fine-grained state-aware graph attention network for dialogue state tracking. *The Journal of Supercomputing*, 81(5):671.
- Weizhe Lin, Bo-Hsiang Tseng, and Bill Byrne. 2021a. [Knowledge-aware graph-enhanced GPT-2 for dialogue state tracking](#). In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, pages 7871–7881, Online and Punta Cana, Dominican Republic. Association for Computational Linguistics.
- Weizhe Lin, Bo-Hsiang Tseng, and Bill Byrne. 2021b. Knowledge-aware graph-enhanced gpt-2 for dialogue state tracking. *arXiv preprint arXiv:2104.04466*.
- George A Miller. 1995. Wordnet: a lexical database for english. *Communications of the ACM*, 38(11):39–41.
- Elliot Nelson, Georgios Kollias, Payel Das, Subhajit Chaudhury, and Soham Dan. 2024. Needle in the haystack for memory based large language models. *arXiv preprint arXiv:2407.01437*.
- Tim Paek and Roberto Pieraccini. 2008. Automating spoken dialogue management design using machine learning: An industry perspective. *Speech communication*, 50(8-9):716–729.
- Abhinav Rastogi, Xiaoxue Zang, Srinivas Sunkara, Raghav Gupta, and Pranav Khaitan. 2020. Towards scalable multi-domain conversational agents: The schema-guided dialogue dataset. In *Proceedings of the AAAI conference on artificial intelligence*, volume 34, pages 8689–8696.
- Amit Singhal. 2012. [Introducing the knowledge graph: things, not strings](#). Google Blog post.
- Zhiqing Sun, Zhi-Hong Deng, Jian-Yun Nie, and Jian Tang. 2019. Rotate: Knowledge graph embedding by relational rotation in complex space. *arXiv preprint arXiv:1902.10197*.
- Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Brian Ichter, Fei Xia, Ed Chi, Quoc Le, and Denny Zhou. 2023. [Chain-of-thought prompting elicits reasoning in large language models](#). *Preprint*, arXiv:2201.11903.
- Jason D Williams, Antoine Raux, and Matthew Henderson. 2016. The dialog state tracking challenge series: A review. *Dialogue & Discourse*, 7(3):4–33.

- Seungpil Won, Heeyoung Kwak, Joongbo Shin, Janghoon Han, and Kyomin Jung. 2023. Break: Breaking the dialogue state tracking barrier with beam search and re-ranking. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 2832–2846.
- Chien-Sheng Wu, Andrea Madotto, Ehsan Hosseini-Asl, Caiming Xiong, Richard Socher, and Pascale Fung. 2019. Transferable multi-domain state generator for task-oriented dialogue systems. *arXiv preprint arXiv:1905.08743*.
- Peng Wu, Bowei Zou, Ridong Jiang, and AiTi Aw. 2020. Gcdst: A graph-based and copy-augmented multi-domain dialogue state tracking. In *Findings of the Association for Computational Linguistics: EMNLP 2020*, pages 1063–1073.
- Shiquan Yang, Rui Zhang, and Sarah Erfani. 2020. Graphdialog: Integrating graph knowledge into end-to-end task-oriented dialogue systems. *arXiv preprint arXiv:2010.01447*.
- Fanghua Ye, Jarana Manotumruksa, and Emine Yilmaz. 2021a. Multiwoz 2.4: A multi-domain task-oriented dialogue dataset with essential annotation corrections to improve state tracking evaluation. *arXiv preprint arXiv:2104.00773*.
- Fanghua Ye, Jarana Manotumruksa, Qiang Zhang, Shenghui Li, and Emine Yilmaz. 2021b. Slot self-attentive dialogue state tracking. In *Proceedings of the web conference 2021*, pages 1598–1608.
- Fanghua Ye, Xi Wang, Jie Huang, Shenghui Li, Samuel Stern, and Emine Yilmaz. 2022. Metaassist: robust dialogue state tracking with meta learning. *arXiv preprint arXiv:2210.12397*.
- Steve Young, Milica Gašić, Blaise Thomson, and Jason D Williams. 2013. Pomdp-based statistical spoken dialog systems: A review. *Proceedings of the IEEE*, 101(5):1160–1179.
- Haoning Zhang, Junwei Bao, Haipeng Sun, Youzheng Wu, Wenye Li, Shuguang Cui, and Xiaodong He. 2022. Monet: Tackle state momentum via noise-enhanced training for dialogue state tracking. *arXiv preprint arXiv:2211.05503*.
- Jeffrey Zhao, Raghav Gupta, Yuan Cao, Dian Yu, Mingqiu Wang, Harrison Lee, Abhinav Rastogi, Izhak Shafran, and Yonghui Wu. 2022. Description-driven task-oriented dialog modeling. *arXiv preprint arXiv:2201.08904*.
- Han Zhou, Ignacio Iacobacci, and Pasquale Minervini. 2022. Xqa-dst: Multi-domain and multilingual dialogue state tracking. *arXiv preprint arXiv:2204.05895*.
- Li Zhou and Kevin Small. 2019a. Multi-domain dialogue state tracking as dynamic knowledge graph enhanced question answering. *arXiv preprint arXiv:1911.06192*.
- Li Zhou and Kevin Small. 2019b. Multi-domain dialogue state tracking as dynamic knowledge graph enhanced question answering. *arXiv preprint arXiv:1911.06192*.
- Meng Zhu and Xiaolong Xu. 2025. Ecdg-dst: A dialogue state tracking model based on efficient context and domain guidance for smart dialogue systems. *Computer Speech & Language*, 90:101741.

A ConvoTree Implementation Details

This appendix provides detailed implementation specifications for ConvoTree, complementing the high-level overview presented in Section 6.

A.1 System Architecture

ConvoTree consists of four interconnected components that operate on different timescales to maintain conversational coherence:

A.1.1 Real-Time Graph Construction

The core API processes each dialogue turn through a multi-stage pipeline. When a new message arrives via the `/api/compress` endpoint, the system employs a two-phase knowledge extraction process:

Phase 1: Entity and Relation Extraction. A lightweight entity and relation extractor identifies salient facts from the current turn, formatting them as RDF triples (*subject, predicate, object*). The extractor uses a combination of named entity recognition (NER) and dependency parsing to identify entities and their relationships within the utterance.

Phase 2: Neural Consolidation. A neural consolidation module merges these facts with the existing conversation graph, resolving entity coreferences and updating relationship weights based on recency and importance. The module employs a graph neural network (GNN) to propagate information across related entities and determine optimal merge strategies.

This dynamic graph construction maintains computational efficiency through incremental updates, avoiding the need to reconstruct the entire graph at each turn.

A.1.2 Neural Memory Consolidation

ConvoTree’s distinctive feature is its neural “dreaming” process that periodically consolidates conversation memory during idle periods. This consolidation engine implements a three-phase cycle inspired by biological memory formation:

Replay Phase: The system reconstructs conversation segments from the knowledge graph, generating synthetic dialogue exchanges that capture key relationship patterns. This phase uses a graph-to-text generation model to create plausible conversation snippets that preserve important factual content while reducing redundancy.

Consolidation Phase: A specialized transformer model processes these reconstructed conversations alongside the original transcripts, identifying stable patterns and compressing redundant information. The model employs attention mechanisms to weight the importance of different conversation elements and determines which information should be preserved versus compressed.

Integration Phase: Consolidated memories are integrated into a hierarchical knowledge representation that balances detail preservation with storage efficiency. Recent information maintains high granularity, while older information is progressively abstracted into higher-level relationship patterns.

A.1.3 LLM-Powered Reasoning System

The reasoning component bridges graph-structured memory with natural language generation through a hybrid architecture:

Retrieval Stage: When processing queries via `/api/chat`, the system first retrieves relevant graph neighborhoods using learned embedding similarity. Graph nodes and edges are embedded using a pre-trained knowledge graph embedding model, enabling efficient similarity-based retrieval.

Integration Stage: Retrieved structural features are combined with compressed conversation summaries and current context. The system uses cross-attention mechanisms to align graph-based information with textual content.

Generation Stage: The integrated representation is fed to a fine-tuned language model that generates contextually appropriate responses. The model architecture includes specialized attention heads for processing graph-structured inputs alongside traditional textual attention.

A.1.4 Long-Term Context Management

ConvoTree maintains conversation continuity across extended time periods through a multi-resolution memory hierarchy:

Short-term Memory (Hours): Recent conversations are stored with full fidelity, including complete dialogue history and detailed graph representations.

Medium-term Memory (Days-Weeks): Interactions are compressed into structured summaries with preserved key relationships. The system maintains a balance between storage efficiency and information retention through adaptive compression rates.

Long-term Memory (Months+): Patterns are distilled into learned relationship embeddings that influence future consolidation decisions without requiring explicit storage of conversation details.

A.2 Implementation Specifications

A.2.1 API Design

ConvoTree exposes a RESTful API with the following core endpoints:

- `POST /api/compress`: Processes new dialogue turns and updates the conversation graph. Accepts JSON payload with message content, speaker identification, and optional metadata.
- `POST /api/chat`: Generates responses based on current conversation state and retrieved memory. Returns both the response text and relevant graph visualizations.
- `GET /api/status`: Provides system status including memory usage, consolidation queue status, and performance metrics.
- `GET /api/graph`: Returns conversation graph in various formats (GraphML, JSON, PNG visualization).

A.2.2 Data Structures

The system employs several key data structures:

Conversation Graph: Implemented as a directed multigraph where nodes represent entities (people, places, concepts) and edges represent relationships with associated confidence scores and temporal metadata.

Memory Hierarchy: Structured as a time-based index with pointers to different resolution levels. Each level maintains appropriate data structures (full transcripts, compressed summaries, or learned embeddings).

Consolidation Queue: Priority queue for managing background consolidation tasks, with priorities based on conversation importance, recency, and available computational resources.

A.2.3 Performance Optimizations

Several optimizations enable production deployment:

Incremental Processing: Graph updates use incremental algorithms to avoid full recomputation. Only affected subgraphs are modified when new information arrives.

Asynchronous Consolidation: Memory consolidation runs asynchronously during conversation pauses, ensuring minimal impact on response latency.

Caching Strategies: Frequently accessed graph neighborhoods and conversation summaries are cached using LRU policies with periodic invalidation.

Resource Management: The system includes adaptive resource allocation based on conversation load and available hardware resources.

A.3 Evaluation Framework

A.3.1 Datasets

ConvoTree evaluation employs both existing benchmarks and custom datasets:

MultiWOZ Extension: Extended conversations created by chaining multiple MultiWOZ dialogues with bridging context to create realistic long-form interactions.

Customer Service Logs: Anonymized customer service conversations spanning multiple sessions, providing realistic examples of long-term context dependencies.

Synthetic Long Conversations: Generated conversations with known ground truth for controlled evaluation of memory retention and consolidation effectiveness.

A.3.2 Metrics

Evaluation focuses on three primary dimensions:

Memory Efficiency: Compression ratio (original size / compressed size), factual accuracy preservation, and retrieval precision for historical information.

Response Quality: Human preference ratings, consistency scores across conversation segments, and factual accuracy of generated responses.

Computational Performance: Response latency, memory usage, consolidation throughput, and scalability metrics across different conversation lengths.

A.4 Future Extensions

The modular architecture supports several planned extensions:

Federated Learning: Extending consolidation across multiple conversation instances while preserving privacy through differential privacy techniques.

Multimodal Integration: Supporting visual and audio inputs through specialized graph node types and cross-modal attention mechanisms.

Domain Adaptation: Automatic adaptation to new conversation domains through meta-learning approaches applied to the consolidation process.

Explainability: Enhanced visualization and explanation capabilities for understanding consolidation decisions and graph evolution patterns.